



# Dynamic Binary Instrumentation

for Security Analysis of Programs



# Meet DBI

**Dynamic Binary Instrumentation** (DBI) is the process of inserting additional code into a running binary to gather execution information and possibly alter its run-time behavior. DBI does not require modifications to the binary itself.

Why DBI for this seminar?

- very popular in PL, systems, and security research
- can help you gain insights in problems you are studying
- you may use it to implement research solutions

# Outline

What we will cover today

1. Introduction to DBI
2. Basic tutorial with Intel Pin
3. Hands-on exercises
4. Usage tips, evasions & wrap-up

# What is DBI?

## Formal definition

*A DBI system is a user-space execution runtime for process-level virtualization.*

An unmodified binary executes within the DBI system as it would happen in a native execution, but introspection and execution control capabilities become accessible to DBI users through well-defined APIs.

The system and the program being analyzed share the same address space.

# Key features

Some notable features:

- no need for source code
- high degree of compatibility
- can expose accurate real-time insights
- has intuitive primitives with flexible granularity (e.g., blocks, functions)
- may handle obfuscated code

DBI can be an ideal solution for analyzing real-world complex software and environments with unpredictable behavior.

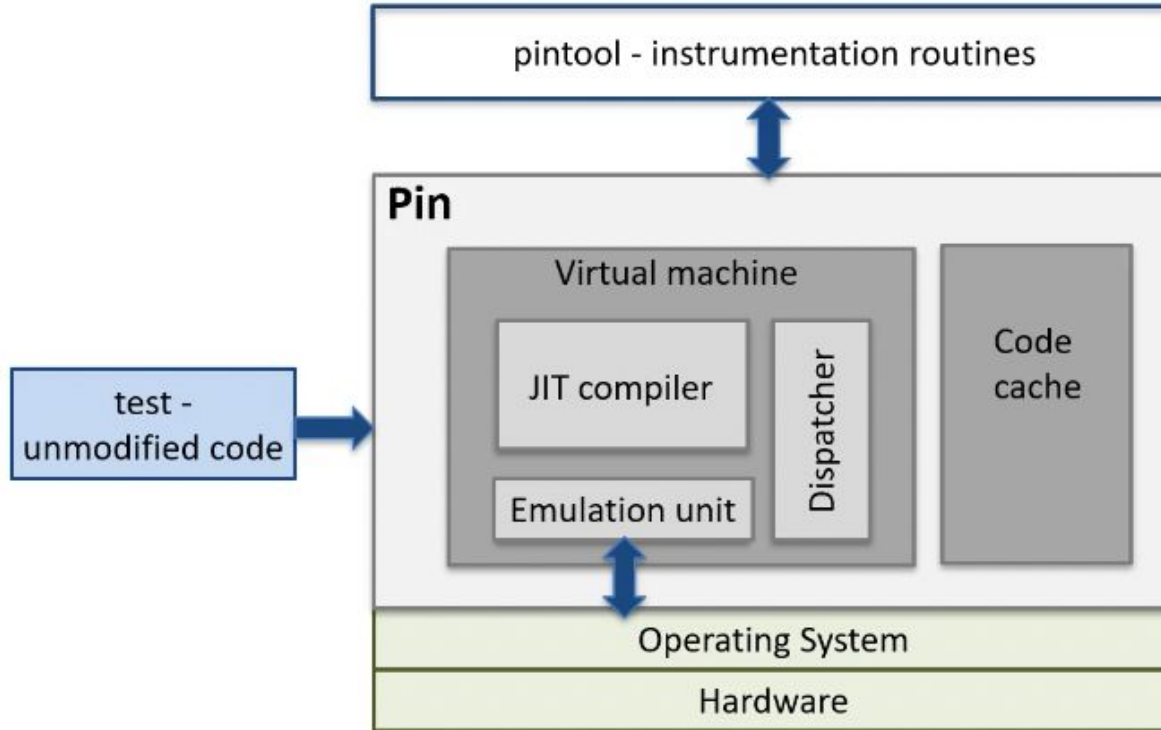
# Workflow

High-level functioning of a DBI system:

1. load the target binary
2. pick some instructions for execution (e.g., until next branch)
3. analyze instructions and apply user-specified instrumentation
4. **just-in-time compile** everything into an executable fragment
5. add the fragment to a code cache and **run it**
6. upon fragment end, go back to step 2

Optimizations (e.g., stitch fragments) mitigate intrinsic overheads

# Workflow



1. load the binary
2. pick instructions
3. instrument them
4. JIT-compile them
5. cache the JIT'ed fragment and run it
6. go back to step 2

# Differences with static instrumentation

Main alternatives: compiler-based instrumentation, static binary rewriting

## Advantages of DBI

- plug-and-play
- support for “dynamic” code
- very flexible instrumentation

## Main drawbacks

- performance overhead
- conspicuous traits



# Differences with static instrumentation

Main alternatives: compiler-based instrumentation, static binary rewriting

## Good use cases for DBI

- profiling behaviors
- debugging & understanding
- hard-to-analyze code

## Good use cases for static techniques

- performance testing
- persistent instrumentation

# Key DBI techniques


Instrumentation can address multiple aspects of execution. Typically:

- **Instruction** instrumentation (by specific types or addresses)
- **Basic block** instrumentation
- **Function** instrumentation (entry/exit events and arguments)
- **Memory access** monitoring

Additional hooks exist for exception handling, process/thread lifecycle, etc.

Users can analyze the entire userland stack or focus only on program code.

# Main DBI systems



FRIDA



*Dyn  
inst*



# Some prominent security uses

SoK work from 2019 reviewed 90+ security papers using DBI for:

- Software protection (30)
- Vulnerability detection (22)
- Malicious software (14)
- Reverse engineering (9)
- Information flow tracking (8)
- Cryptoanalysis (7)

and more. DBI tools can also help with feasibility studies for larger projects.

# Writing a DBI tool with Pin

Today we focus on Intel Pin for its broad capabilities and compatibility.

Suggested steps

1. Identify what events you want to monitor/alter
2. Sketch callbacks for analyzing event instances
3. Set up a Pin tool skeleton
4. Register such callbacks with appropriate instrumentation
5. Try them live and continue expanding your tool

# Typical instrumentation tasks

## Profiling

- Count instructions/blocks/calls/accesses
- Identify targets (e.g., indirect calls)
- Measure execution time
- Log data (e.g., API parameters)

## Debugging/understanding

- Validate memory accesses
- Intercept errors
- Track execution context (call stack)

# Getting started with Pin

```
1  #include "pin.H"
2
3  int main(int argc, char *argv[]) {
4      // Initialize symbol processing
5      PIN_InitSymbols();
6
7      // Initialize Pin
8      PIN_Init(argc, argv);
9
10     // Add instrumentation function for instructions
11     INS_AddInstrumentFunction(InstructionCallback, NULL);
12
13     // Start the program execution
14     PIN_StartProgram();
15
16     return 0;
17 }
```

To simplify the build system, borrow the Makefile rules from the SimpleExamples folder of all Pin releases.

# Hands-on



<https://github.com/dcdelia/cuso24240006>

All you need is a Linux system with g++,  
build-essentials, and Intel Pin.

Setup instructions are in the GitHub page.

(short link: <https://tinyurl.com/cusodbi>)



# Building our Pin tool #1

We start with a classic task: tracing and counting all executed instructions.

Learning goals

- Instrumentation-time vs analysis-time callbacks
- INS instrumentation: `INS_InsertCall` with `IPOINT_BEFORE` ([docs](#))
- Handling execution termination

# Building our Pin tool #2

We now move to a classic security scenario: intercepting indirect calls.

Learning goals

- Distinguishing instruction types: [INS\\_IsIndirectControlFlow](#), [INS\\_IsCall](#)
- Inspecting [run-time data](#): IARG\_INST\_PTR, IARG\_BRANCH\_TARGET\_ADDR
- IMG instrumentation for whole modules
- RTN abstraction for looking up functions ([docs](#))

# Building our Pin tool #3

We will maintain a stack trace and present it to the user at a crash site.

Learning goals

- More instruction types ([INS\\_IsDirectCall](#), [INS\\_IsRet](#))
- More run-time data for callbacks
- Intercepting signals: [PIN\\_AddContextChangeFunction](#) and [callback](#)
- Validating consistency of retrieved run-time data
- Accessing debug information

# Implementation tricks from experience

DBI can be **very expensive**. Design work can mitigate these costs:

- Conducting work at instrumentation (vs. analysis) time is preferable
- Simple analysis callbacks may be inlined
- Try to trade space for speed when possible
- Advanced means (e.g., predicated instrumentation) for intensive tracing

# Conspicuous features of DBI



**SoK: Using DBI for Security (And How You May Get Caught Red Handed) - ASIACCS'19**

# What does transparency mean?

For **DBI architects**: support program semantics as in a native execution.

*«The further we push transparency, the more difficult it is to implement, while at the same time fewer applications require it.» - DynamoRIO's creator*

For **security researchers**: possibility of adversarial sequences.



# Opportunities for attackers

The ASIACCS'19 SoK identifies 7 attack surfaces:

- time overhead
- memory contents and permissions
- leaking a code cache address
- DBI engine internals
- interactions with OS
- exception handling
- translation defects

# Examples of detections


## > Time overhead

- *general slowdown from dynamic analysis*
- *ad-hoc detections*



## > Memory contents and permissions

- *extra sections, data patterns*
- *increased memory usage*
- *stress layout for **consistency***  
*(e.g. NX policy, guard pages)*



## > Leaking an address from code cache



# How to counter DBI evasions

*Theoretically weak, but practically effective: **patch** obvious imperfections.*

Domain experts are aware of the most recurrent possibilities. For example, for malware analysis, we can patch Pin's imperfect handling of `int 2d`.

Some cases are hard: one can then use DBI to conduct a **feasibility** study and then devise later a more complex implementation with other techniques.

# What affects a researcher's choice?

Some relevant factors:

- required instrumentation capabilities
- altering execution can be complex with other technologies
- deployability of the runtime
- tricky behaviors (e.g., sensitivity to time variations)
- characteristics of the adversary

Thanks!